



# Administration & Operations

Planning / Testing

◁EROSPIKE▷

# Objectives

At the end of this module, you will be able to:

- Have a basic sense of capacity planning.
- Load test data into an Aerospike database using the Aerospike Java Benchmark tool.
- Run benchmark tests.



# Capacity Planning

# Capacity Planning

For a exact, detailed planning guide, go to:

<http://www.aerospike.com/docs/operations/plan/capacity/>

For a quick estimate (no secondary indexes) you will need the following:

- number of records
- replication factor
- size of records
- number of bins
- average set size
- type of storage (RAM or flash/SSD)

# Capacity Planning – Quick Estimate

Area	How stored	Formula	Note
Primary Index	RAM	$n * r * 64$	The amount of RAM needed for the primary index is fixed at 64 bytes.
Data storage	RAM	$n * r * (2 + (17 * b) + v)$	Every objects needs 2 bytes for overhead, 17 bytes per bin, and the actual data
Data storage	Flash/SSD	$n * r * p$  Where p is $(64 + (9 + s) + (28 * b) + 5 + v)$ -> round up to nearest 128 bytes	Every object needs to store the index (64 bytes), set overhead (9 +s bytes) , general overhead (28 bytes), type info (5 bytes), and the actual data. Because Aerospike stores data in 128 byte blocks, you must round up to the nearest 128 byte amount.

- $n$  = number of records
- $r$  = replication factor
- $v$  = average size of records
- $b$  = number of bins
- $s$  = average set size



## Benchmarks

# Benchmarking

## Goal: **Saturate the cluster**

- Start with single server – BEWARE: “writes” values are not representative
- Change **1 thing** at a time
  - Move toward production configuration
- Are you **bottlenecked** on the
  - Client
  - Network

## Note:

- Don't expect high performance with **VMs**. You must be very careful when testing in a public cloud such as Amazon.
- Don't expect realistic number with client and server on the **same host**.

# The Java Benchmark Tool

- Multi-platform
  - Portable
  - Easy to use
  - Total throughput easily seen from database nodes
- 
- Poor performance with objects larger than 4 KB
  - Run multiple instances on multiple servers to get to high throughput



# Java Benchmark Syntax

```
> ./run_benchmarks [options]
```

Flag	Description
<code>-b, --bins &lt;arg&gt;</code>	The number of bins (like columns) of data.
<code>-g, --throughput &lt;arg&gt;</code>	The target total throughput. This will include all operations.
<code>-h, --host &lt;arg&gt;</code>	A seed node (any node in the cluster). The tool will learn about the other nodes from this one.
<code>-p, --port &lt;arg&gt;</code>	The port on the Aerospike node to connect to (default 3000)
<code>-k, --keys &lt;arg&gt;</code>	This is the number of keys/records to use. If you select a number too small, you will get write contention. Too large and you may run out of memory.
<code>-s, --set &lt;arg&gt;</code>	The set where the data will be loaded.
<code>-latency &lt;arg&gt;</code>	Produce a table of latencies as measured from the client. This can be compared with the latencies on the server to see where any bottleneck may lie. <b>Recommended value “-latency 7,1”</b>
<code>-n, --namespace &lt;arg&gt;</code>	The namespace to test against. This is similar to a “database” in a relational database.
<code>-o, --objectSpec &lt;arg&gt;</code>	The type of object to use: “I” is an integer, “S:<size>” is the size of the object. Is is recommended to start with small objects to remove the network
<code>-w, --workload &lt;arg&gt;</code>	Type of workload: I - Insert only. test stops when the final record is loaded (based on –k above) RU,<percent> - The read percentage. “90” means a 90% read rate.
<code>-z, --threads &lt;arg&gt;</code>	The number of threads used by the benchmark. For bare metal, 64 is a good value. For VMs, you should use 4-8 threads/core.

# The Java Benchmark Tool

The Aerospike benchmark tool can be found within the full Aerospike Java Client.

Untar, go to the appropriate directory and compile it:

- `cd /home/aerotrain/pkgs/aerospike`
- `tar xvf aerospike-client-java.tgz`
- `cd aerospike-client-java-*/benchmarks`
- `mvn package`



## **Loading Test Data**

# Loading Data Example

Prior to running any tests, you should load data into the database.

```
./run_benchmarks    -h 127.0.0.1 \    # Local DB server
                    -p 3000 \    # Local port
                    -n test \    # Namespace "test"
                    -k 100000 \   # Load 100,000 records
                    -s testset \   # Use the set "testset"
                    -latency "7,1" \ # Show latency numbers
                    -o S:100 \    # Use 100 byte string
                    -w I \    # Insert only
                    -z 8          # Use 8 threads
```

➤ `./run_benchmarks -h 127.0.0.1 -p 3000 -n test -s testset -k 100000 -latency "7,1" -S 1 -o S:100 -w I -z 8`

When the script has loaded all the data, it will exit.

# Loading Data Example Output

```
$/run_benchmarks -h 127.0.0.1 -p 3000 -n test -s testset -k 100000 -latency "7,1" -S 1 -o  
S:100 -w I -z 8
```

```
Benchmark: 127.0.0.1:3000, namespace: test, set: testset, threads: 8, workload: INITIALIZE  
keys: 100000, start key: 1, key length: 10, bins: 1, throughput: unlimited, debug: false  
write policy: timeout: 0, maxRetries: 2, sleepBetweenRetries: 500  
bin type 1: string[100]  
2014-08-27 22:12:52.091 INFO Thread 1 Add node BB9ED14150B0022 127.0.0.1:3000  
2014-08-27 22:12:52.137 write(count=0 tps=0 timeouts=0 errors=0)  
    <=1ms >1ms >2ms >4ms >8ms >16ms >32ms  
write    0%    0%    0%    0%    0%    0%    0%  
2014-08-27 22:12:53.141 write(count=2432 tps=2432 timeouts=0 errors=0)  
    <=1ms >1ms >2ms >4ms >8ms >16ms >32ms  
write   90%   10%    9%    8%    7%    5%    1%  
2014-08-27 22:12:54.170 write(count=6492 tps=4060 timeouts=0 errors=0)  
    <=1ms >1ms >2ms >4ms >8ms >16ms >32ms  
write   92%    8%    7%    6%    5%    3%    1%  
2014-08-27 22:12:55.194 write(count=12705 tps=6213 timeouts=0 errors=0)  
...
```

**While this is running, check on the AMC.**

# Latency measurement

Latency parameter determines how latencies are displayed.

Format:

“-latency N,E[,us]”

N is the number of buckets

E is the number of powers of 2 between each bucket

us to get microsecond resolution on timings (Default is ms)

Buckets show percent transaction greater than that number (except 1<sup>st</sup> one)

-latency “7,1”

<=1ms	>1ms	>2ms	>4ms	>8ms	>16ms	>32ms	>64ms	>128ms	...
58%	42%	27%	15%	3%	1%	0%	0%	0%	

-latency “5,2”

<=1ms	>1ms	>2ms	>4ms	>8ms	>16ms	>32ms	>64ms	>128ms	...
58%	42%	27%	15%	3%	1%	0%	0%	0%	



## **Running Benchmarks**

# Running Balanced Workload

The first test to run is on a balanced workload. This will test database operations that are **50% reads 50% writes** (and updates).

```
./run_benchmarks    -h 127.0.0.1 \    # Local DB server
                    -p 3000 \    # Local port
                    -n test \    # Namespace "test"
                    -k 100000 \    # Load 100,000 records
                    -s testset \    # Use the set "testset"
                    -latency "7,1" \    # Show latency numbers
                    -o S:100 \    # Use 100 byte string
                    -w RU,50 \    # Read at 50%
                    -z 8          # Use 8 threads
```

This script will continue running until it is stopped.



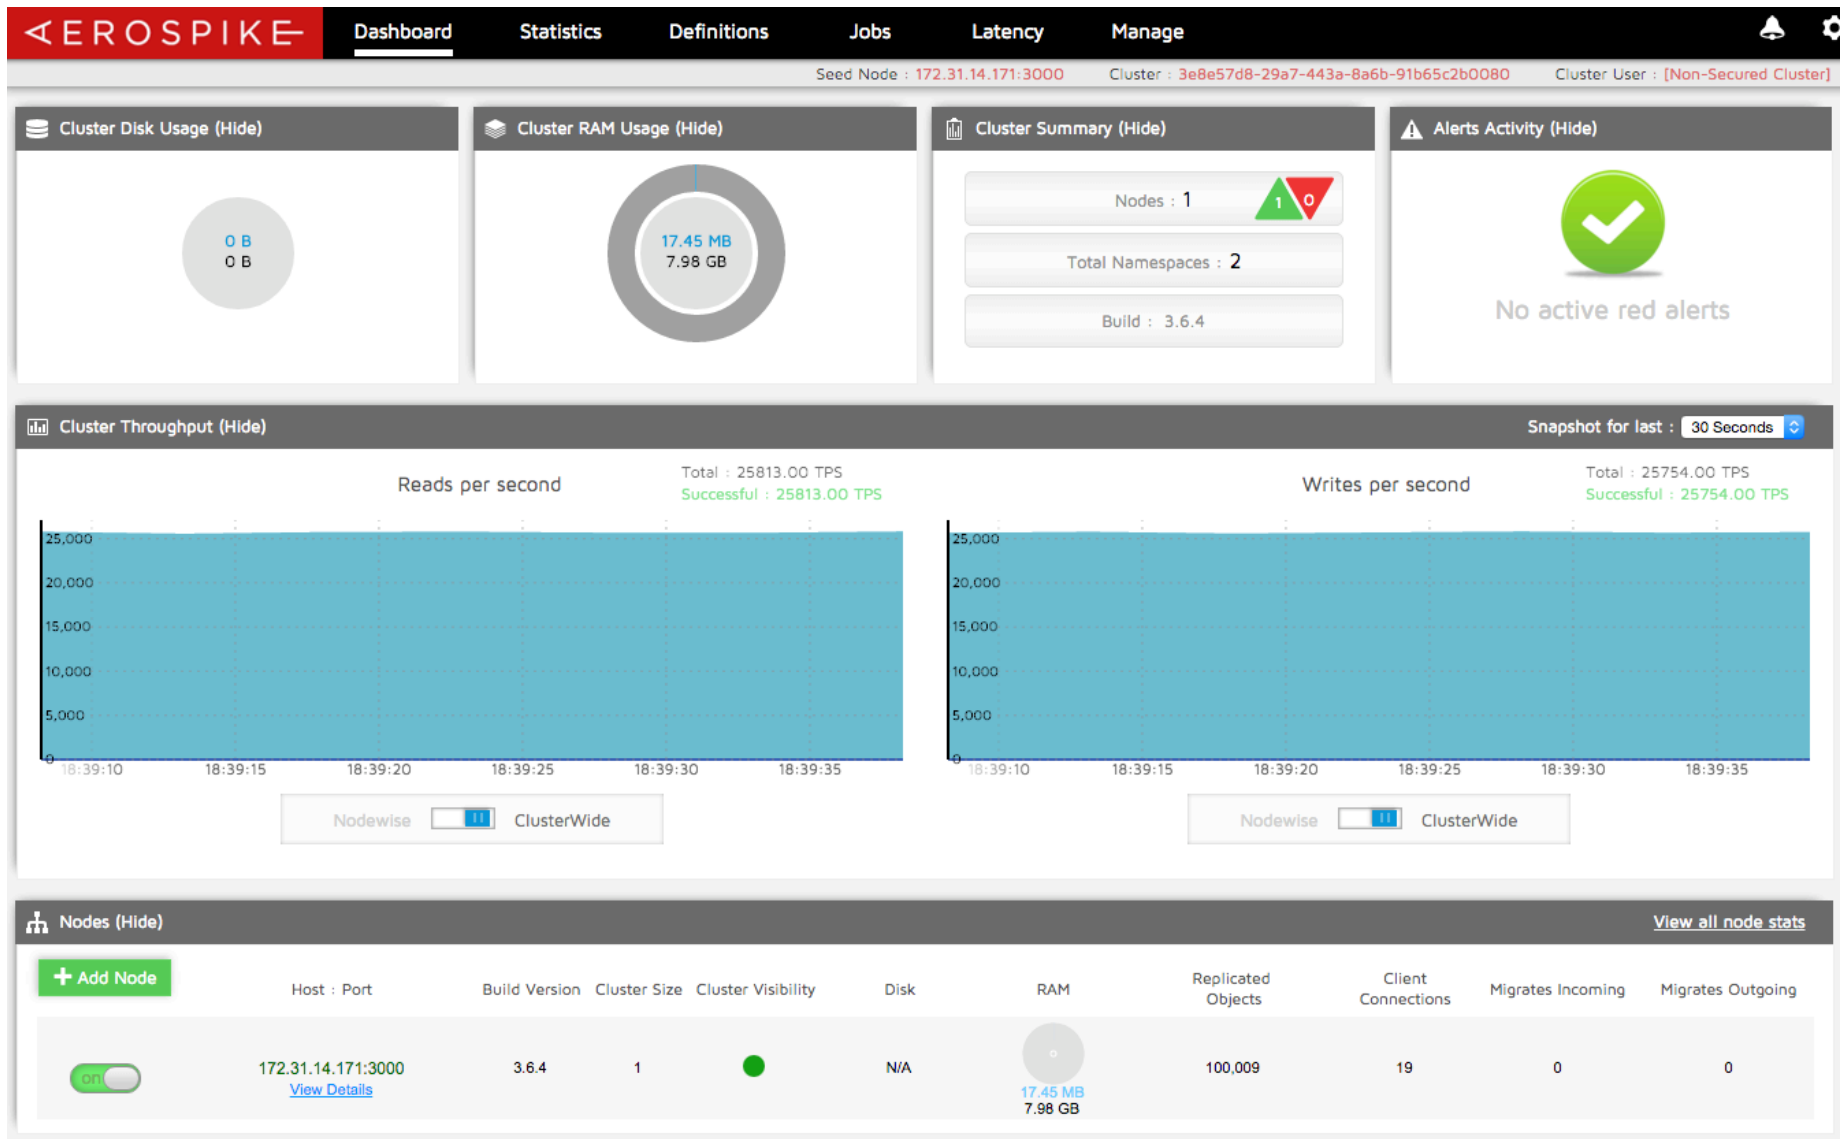
# Running Balanced Load Example Output

```
$/run_benchmarks -h 127.0.0.1 -p 3000 -n test -s testset -k 100000 -latency "7,1" -o S:100 -w RU,50 -z 8
```

```
Benchmark: 127.0.0.1:3000, namespace: test, set: testset, threads: 8, workload: READ_UPDATE
read: 50% (all bins: 100%, single bin: 0%), write: 50% (all bins: 100%, single bin: 0%)
keys: 100000, start key: 0, key length: 10, bins: 1, throughput: unlimited, debug: false
read policy: timeout: 0, maxRetries: 2, sleepBetweenRetries: 500, reportNotFound: false
write policy: timeout: 0, maxRetries: 2, sleepBetweenRetries: 500
bin type 1: string[100]
2014-08-27 22:14:56.269 INFO Thread 1 Add node BB9ED14150B0022 127.0.0.1:3000
2014-08-27 22:14:56.329 write(tps=2 timeouts=0 errors=0) read(tps=0 timeouts=0 errors=0)
total(tps=2 timeouts=0 errors=0)
    <=1ms >1ms >2ms >4ms >8ms >16ms >32ms
write    89%  11%  11%  11%  11%    5%    2%
read     86%  14%  13%  13%  13%    9%    2%
2014-08-27 22:14:57.577 write(tps=1458 timeouts=0 errors=0) read(tps=1402 timeouts=0 errors=0)
total(tps=2860 timeouts=0 errors=0)
    <=1ms >1ms >2ms >4ms >8ms >16ms >32ms
write    86%  14%  13%  12%  11%    6%    1%
read     89%  11%  10%   9%   8%    5%    2%
2014-08-27 22:14:58.585 write(tps=1994 timeouts=0 errors=0) read(tps=1934 timeouts=0 errors=0)
total(tps=3928 timeouts=0 errors=0)
    <=1ms >1ms >2ms >4ms >8ms >16ms >32ms
write    93%   7%   7%   6%   5%    3%    1%
read     94%   6%   5%   5%   4%    2%    1%
...
```

**While this is running, check on the AMC.**

# AMC running benchmark



# Running High-read Workload

You can then run a workload based on high read rates. This will test database operations that are **95% reads 5% writes**

```
./run_benchmarks    -h 127.0.0.1 \    # Local DB server
                    -p 3000 \    # Local port
                    -n test \    # Namespace "test"
                    -k 100000 \  # Load 100,000 records
                    -s testset \  # Use the set "testset"
                    -latency "7,1" \ # Show latency numbers
                    -o S:100 \    # Use 100 byte string
                    -w RU,95 \  # Read at 95%
                    -z 8          # Use 8 threads
```

This script will continue running until it is stopped.

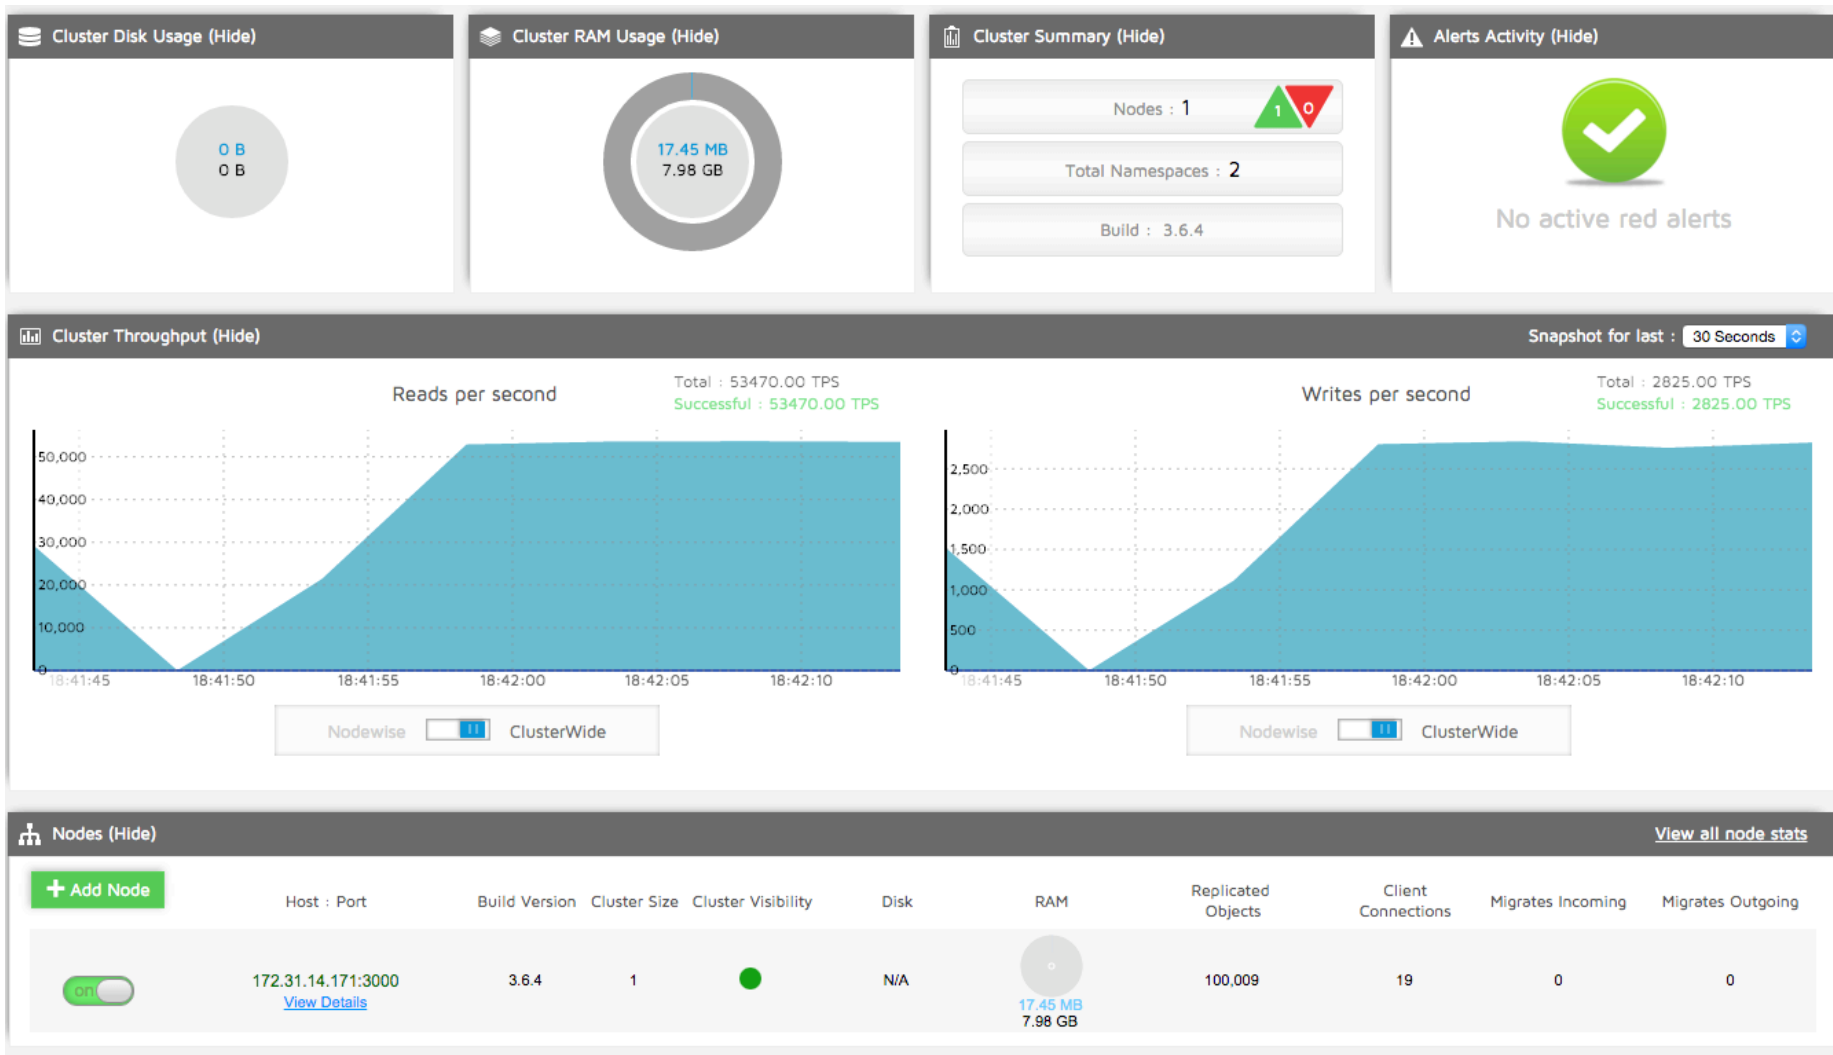
# Running High-read workload Example Output

```
$. /run_benchmarks -h 127.0.0.1 -p 3000 -n test -s testset -k 100000 -latency "7,1" -o S:100 -w RU,95 -z 8
```

```
Benchmark: 127.0.0.1:3000, namespace: test, set: <empty>, threads: 8, workload: READ_UPDATE
read: 95% (all bins: 100%, single bin: 0%), write: 5% (all bins: 100%, single bin: 0%)
keys: 100000, start key: 0, key length: 10, bins: 1, throughput: unlimited, debug: false
read policy: timeout: 0, maxRetries: 2, sleepBetweenRetries: 500, reportNotFound: false
write policy: timeout: 0, maxRetries: 2, sleepBetweenRetries: 500
bin type 1: string[100]
2015-05-28 18:58:09.946 INFO Thread 1 Add node BB94FB6A4647106 127.0.0.1:3000
...
2015-05-28 18:58:17.356 write(tps=488 timeouts=0 errors=0) read(tps=8374 timeouts=0 errors=0)
total(tps=8862 timeouts=0 errors=0)
      <=1ms >1ms >2ms >4ms >8ms >16ms >32ms
write   96%   4%   3%   2%   2%   1%   1%
read    97%   3%   3%   2%   2%   1%   1%
2015-05-28 18:58:18.391 write(tps=468 timeouts=0 errors=0) read(tps=9569 timeouts=0 errors=0)
total(tps=10037 timeouts=0 errors=0)
      <=1ms >1ms >2ms >4ms >8ms >16ms >32ms
write   95%   5%   4%   3%   2%   1%   0%
read    97%   3%   2%   2%   2%   1%   0%
...
```

**While this is running, check on the AMC.**

# AMC running benchmark



# Summary

What we have covered:

- Have a basic sense of capacity planning.
- Load test data into an Aerospike database using the Aerospike Java Benchmark tool.
- Run benchmark tests.