

AEROSPIKE

AS101 Lab Exercises



Lab: Key-value Operations

Objective

After successful completion of this Lab module you will have:

- Connected to a cluster
- Write and read records using simple and complex values
- Used advanced key-value techniques

Lab Overview

The lab exercises add functionality to a simple Twitter-like console application (tweetaspike) using Aerospike as the database.

In this Lab, we will focus on key-value operations and techniques. You will add code to:

- Create users and Tweets
- Read all the Tweets for a user
- Use an advance Key-value feature

The exercises shell is located in your cloned GitHub directory
`~/exercises/Key-valueOperations/<language>`

Make sure you have your server up and you know its [IP address](#)

In your cloned or downloaded repository, you will find the following directories:

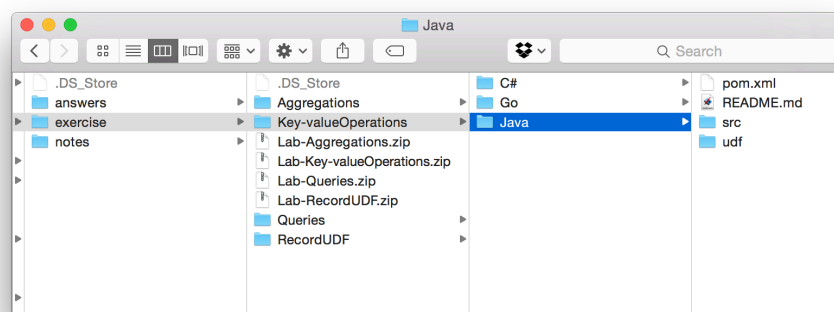
- Answers
- Exercise
- Notes

In the exercise directory, select the subdirectory for your programming language:

- C#
- Java
- Go
- PHP
- Ruby
- Node.js
- Python

The exercises for this module are in the Key-valueOperations directory and you will find a Project/Solution/Codebase that is partly complete. Your tasks is to complete the code as outlined in each exercise.

Make sure you have your server up and you know its IP address



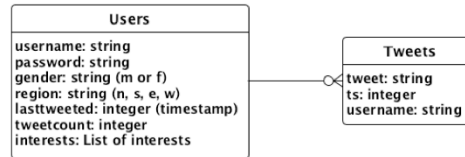
Tweetaspike Data Model

Users

- Namespace: test, Set: users, Key: <username>
- Bins:
 - username - String
 - password - String (For simplicity password is stored in plain-text)
 - gender - String (Valid values are 'm' or 'f')
 - region - String (Valid values are: 'n' (North), 's' (South), 'e' (East), 'w' (West) -- to keep data entry to minimal we just store the first letter)
 - lasttweeted - int (Stores epoch timestamp of the last/most recent tweet) -- Default to 0
 - tweetcount - int (Stores total number of tweets for the user) -- Default to 0
 - interests - Array of interests

Tweets

- Namespace: test, Set: tweets, Key: <username:<counter>>
- Bins:
 - tweet - string
 - ts - int (Stores epoch timestamp of the tweet)
 - username - string



Users

Namespace: test, Set: users, Key: <username>

Bin name	Type	Comment
username	String	
password	String	For simplicity password is stored in plain text
region	String	Valid values are: 'n' (North), 's' (South), 'e' (East), 'w' (West) to keep data entry to minimal we just store the first letter
lasttweeted	Integer	Stores epoch timestamp of the last/most recent tweet Default to 0
tweetcount	Integer	Stores total number of tweets for the user – Default 0
Interests	List	A list of interests

Tweets

Namespace: test, Set: tweets, Key: <username:<counter>>

Bin name	Type	Comment
tweet	String	Tweet text
ts	Integer	Stores epoch timestamp of the tweet
username	String	User name of the tweeter

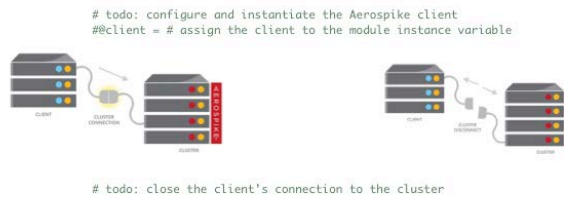


Ruby Exercises

Exercise 1 – Ruby: Connect & Disconnect

Locate the Training module in the Ruby project

1. Create an instance of the Aerospike client in the Training module's init method. Ensure that this connection is created only once
2. Add code to disconnect from the cluster in the Training.finish method. Ensure that this code is executed only once



In this exercise you will connect to a Cluster by creating an Aerospike client instance, using a single IP address and port. These should be to a node in your own cluster.

Ensure that you only create one client instance at the start of the program. The Aerospike is thread safe and creates a pool of worker threads, this means you DO NOT need to create your own connection or thread pool.

1. In the Training.init, add code similar to this;

```
@client = host ? Client.new(Host.new(host, port)) : Client.new
```

Make sure you have your server up and you know its IP address

2. In Training.finish, add code, similar to this, to disconnect from the cluster. This should only be done once. After close() is called, the client instance cannot be used.

```
@client->close
```

Exercise 2 – Ruby: Write User Record

Create a User Record

Locate the UserService module in the Ruby project

1. Create a User Record – In UserService.create_user
 1. Fill the Training.get_user_key method
 2. Use the method to get a key for the new user record
 3. Write the user record

Create a User Record. In UserService.create_user, add code similar to this:

1. Fill the Training.get_user_key method
2. Use the method to get a key for the new user record
3. Write the user record

```
def get_user_key(username)
  Key.new(self.namespace, self.set_name, username)
end

def create_user(client)
  puts "\nCreate a new user".colorize(:color => :blue, :mode
=> :bold)
  print 'Enter username (or hit Return to skip):'
  ".colorize(:blue)
  username = gets.chomp
  bins = { 'username' => username }
  return if username.length < 1
  print "Enter password for #{username}: ".colorize(:blue)
  bins['password'] = gets.chomp
  print "Select gender (f or m) for #{username}:"
  ".colorize(:blue)
  bins['gender'] = gets.chomp
  print "Select region (north, south, east or west) for
#{username}: ".colorize(:blue)
  bins['region'] = gets.chomp
  print "Enter comma-separated interests for #{username}:"
  ".colorize(:blue)
  bins['interests'] = gets.chomp.split(',')

  print "Creating user record > ".colorize(:color
=> :black, :mode => :bold)
  begin
    key = Key.new(self.namespace, self.set_name, username)
    client.put(key, bins, self.write_policy)
    self.yep
  rescue
    self.nope
    puts "Connection to Aerospike cluster failed! Please
check the server settings and try again!".colorize(:color
=> :red, :mode => :bold)
  end
end
```


Exercise 2 – Ruby: Write Tweet Record

Create a Tweet Record

Locate the TweetService module in the Ruby project

1. Create a Tweet Record – In TweetService.create_tweet
 1. Get the user's tweetcount value
 2. Fill the Training.get_tweet_key method
 3. Create the tweet record
 4. Update the user record's tweet count and last tweeted timestamp

Create a Tweet Record. In TweetService.create_tweet, add code similar to this:

1. Get the user's tweetcount value
2. Fill the Training.get_tweet_key method
3. Create the tweet record
4. Update the user record's tweet count and last tweeted timestamp

```
def get_tweet_key(username, id)
  Key.new(self.namespace, 'tweets', "#{username}#{id}")
end

def create_tweet(client, username = nil)
  puts "\nCreate a new tweet".colorize(:color => :blue, :mode
=> :bold)
  unless username
    print "Enter username (or hit Return to skip): ".colorize(:blue)
    username = gets.chomp
  end
  unless username == ''
    key = self.get_user_key(username)
    rec = client.get(key, ['tweetcount'])
    bins = rec.bins unless rec.nil?
    tweet_count = bins.nil? ? 1 : bins['tweetcount'] + 1
  end
  ts = (Time.now.to_f * 1000).round
  bins = {'ts' => ts}
  print "Enter tweet for #{username}: ".colorize(:blue)
  bins['tweet'] = gets.chomp
  bins['username'] = username
  print "Creating tweet record >".colorize(:color => :black, :mode
=> :bold)
  key = self.get_tweet_key(username, tweet_count)
  begin
    client.put(key, bins)
    self.yep
  rescue Exception => e
    self.nope
    puts "Failed to create the tweet".colorize(:color => :red, :mode
=> :bold)
    pp e
    return
  end
  key = self.get_user_key(username)
  bins = {'tweetcount' => tweet_count, 'lasttweeted' => ts}
  print "Updating the user record >".colorize(:color => :black, :mode
=> :bold)
  begin
    client.put(key, bins)
    self.yep
  rescue
    self.nope
    print "Failed to update the user record >".colorize(:color
=> :black, :mode => :bold)
  end
end
```

Exercise 2 Ruby: Read Records

Read User Record

Locate the UserService module in the Ruby project

1. Read a User Record – In UserService.get_user
 1. Read the user record from the cluster

In UserService.get_user, add code, similar to this, to:

```
def get_user(client)
  puts "\nCreate a user".colorize(:color
=> :blue, :mode => :bold)      print 'Enter username
(or hit Return to skip): '.colorize(:blue)
  username = gets.chomp
  key = self.get_user_key(username)
  rec = client.get(key)
  if rec
    puts "Record.key (Key)".colorize(:mode
=> :bold)
    pp rec.key
    puts "Record.generation
(Fixnum)".colorize(:mode => :bold)      pp
rec.generation
    puts "Record.expiration
(Fixnum)".colorize(:mode => :bold)      pp
rec.expiration
    puts "Record.bins (Hash)".colorize(:mode
=> :bold)
  else
    puts "There is no such user
#{username}".colorize(:red)      end
  end
end
```

Exercise 3 – Ruby: Batch Read

Batch Read Tweets for a given user

Locate the TweetService module in the Ruby project

1. In TweetService.batch_get_tweets
 1. Get the value of the user's tweetcount bin
 2. Determine how many tweets the user has
 3. Create an array of tweet key instances
 4. Initiate a batch-read operation
 5. Output the tweets to the console

Read all the tweets for a given user. In TweetService.batch_get_tweets, add code similar to this:

- | | |
|--|---|
| <ol style="list-style-type: none">1. Get the value of the user's tweetcount bin2. Determine how many tweets the user has3. Create an array of tweet key instances4. Initiate a batch-read operation5. Output the tweets to the console | <pre>def batch_get_tweets(client, username = nil) puts "\nGet the user's tweet".colorize(:color => :blue, :mode => :bold) unless username print "Enter username (or hit Return to skip): ".colorize(:blue) username = gets.chomp end unless username == '' key = self.get_user_key(username) rec = client.get(key, ['tweetcount']) unless rec puts "There is no such user #{username}".colorize(:red) return end bins = rec.bins tweet_count = bins.nil? ? 0 : bins['tweetcount'] keys = [] (1..tweet_count).each do i keys.push(self.get_tweet_key(username, i)) end print "Batch-reading the user's tweets >".colorize(:color => :black, :mode => :bold) begin recs = client.batch_get(keys) self.yep rescue Exception => e self.nope puts "Failed to batch-read the tweets for #{username}".colorize(:color => :red, :mode => :bold) pp e return end puts "Here are #{username}'s tweets:".colorize(:color => :blue, :mode => :bold) recs.each do r puts r.bins['tweet'] end end end</pre> |
|--|---|

Exercise 4 – Ruby: Scan

Scan all tweets for all users

Locate the TweetService module in the Ruby project

1. In TweetService.scan_tweets
 1. Initiate scan operation on the test.tweets set
 2. Output the tweets to the terminal

Scan all the tweets for all users – warning – there could be a large result set.

In the TweetService.scan_tweets, add code similar to this:

1. Initiate scan operation on the test.tweets set
2. Output the tweets to the terminal

```
def scan_tweets(client)
  puts "\nScan for tweets".colorize(:color
=> :blue, :mode => :bold)
  begin
    policy = ScanPolicy.new(:fail_on_cluster_change =>
true)
    recordset = client.scan_all(self.namespace,
'tweets', [], policy)

    recordset.each do |rec|
      puts rec.bins['tweet']
    end
  rescue
    puts "Failed to scan test.tweets".colorize(:color
=> :red, :mode => :bold)
  end
end
```

Exercise 5 – Ruby: Read-modify-write

Update the User record with a new password ONLY if the user record is unmodified.

Locate the UserService module in the Ruby project

1. In `UserService.check_and_set_password`
 1. Get the generation of the user record
 2. Instantiate a `WritePolicy`
 3. Set the `WritePolicy.generation` to the value read from the user record
 4. Set `WritePolicy's generation_policy` to `GenerationPolicy::EXPECT_GEN_EQUAL`
 5. Update the user record with the new password

Update the User record with a new password ONLY if the User record is unmodified In `UserService.check_and_set_password`, add code similar to this:

- | | |
|--|---|
| 1. Get the generation of the user record | <code>key = self.get_user_key(username)</code> |
| 2. Instantiate a <code>WritePolicy</code> | <code>rec = client.get_header(key)</code> |
| 3. Set the <code>WritePolicy.generation</code> to the value read from the user record | <code>generation = rec.generation</code> |
| 4. Set <code>WritePolicy's generation_policy</code> to <code>GenerationPolicy::EXPECT_GEN_EQUAL</code> | <code>write_policy = WritePolicy.new</code> |
| 5. Update the user record with the new password | <code>write_policy.generation_policy =</code> |
| | <code>GenerationPolicy::EXPECT_GEN_EQUAL</code> |
| | <code>write_policy.generation = generation</code> |
| | <code>client.put(key, bins, write_policy)</code> |

Exercise 6 – Ruby: Operate

Update the tweet count and timestamp and examine the new tweet count

Locate the TweetService module in the Ruby project

1. In `TweetService.update_tweet_count`
 1. Use the `operate` method to update the user record, passing in policy, user record key, `.add` operation incrementing tweet count, `.put` operation updating timestamp and `.get` operation to read the user record
2. In `TweetService.create_tweet`
 1. Remove the code added in Exercise 2 for updating tweet count and timestamp
 2. Output the updated tweet count to the terminal

In `TweetService.update_tweet_count`

1. Use the `operate` method to update the user record, passing in policy, user record key, `.add` operation incrementing tweet count, `.put` operation updating timestamp and `.get` operation to read the user record

```
ops = [  
  Operation.add(Bin.new('tweetcount', 1)),  
  Operation.put(Bin.new('lasttweeted', (Time.now.to_f *  
    1000).round))  
]  
client.operate(key, ops)
```

In `TweetService.create_tweet`

1. Remove the code added in Exercise 2 for updating tweet count and timestamp
2. Output the updated tweet count to the terminal

Summary

You have learned how to:

- Connect to Cluster
- Write and Read Records
- Batch Read Records
- Read-Modify-Write
- Operate
- Handle errors correctly



Lab: User Defined Functions - record

Objective

After successful completion of this Lab module you will have:

- Coded a Record UDF
- Registered the UDF with a cluster
- Invoked the UDF from your C#, Go, PHP, Ruby, Node.js or Java application

Lab Overview

The lab exercise augments "tweetaspike" by using a Record UDF. Here we will focus on a Record UDF that updates user password.

You will:

- Write a user defined function, in Lua, to update the user password
- Register the UDF
- Execute the UDF from your application

The application shell is located in your cloned GitHub directory

`~/exercises/RecordUDF/<language>`

Make sure you have your server up and you know its [IP address](#)

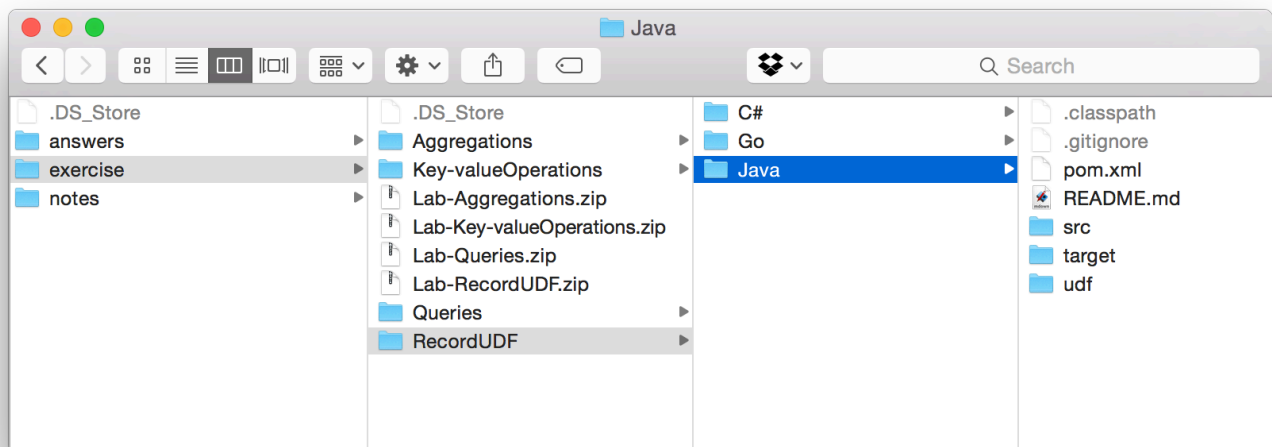
In your cloned or downloaded repository you will find the following directories:

- Answers
- Exercise
- Notes

In the exercise directory, select the subdirectory for your programming language:

- C#
- Java
- Node.js
- Go
- PHP
- Ruby

The exercises for this module are in the UDF directory and you will find a Project/SoluNon/Codebase that is partly complete. Your task is to complete the code as outlined in each exercise.



Exercise 1 – All languages: Write Record UDF

Locate updateUserPwd.lua file in the **udf** folder

1. Log current password
2. Assign new password to the user record
3. Update user record
4. Log new password
5. Return new password

```
function updatePassword(topRec,pwd)
-- Exercise 1
-- TODO: Log current password
-- TODO: Assign new password to the user record
-- TODO: Update user record
-- TODO: Log new password
-- TODO: return new password
end
```

In this exercise you will create a record UDF that:

1. Logs the current password
2. Assigns a new password to the record, passed in via the **pwd** parameter
3. Updates the user record by calling **aerospike:update(topRec)**
4. Logs the new password
5. Returns the new password to the client

```
function updatePassword(topRec,pwd)
-- Log current password
debug("current password: " .. topRec['password'])
-- Assign new password to the user record
topRec['password'] = pwd
-- Update user record
aerospike:update(topRec)
-- Log new password
debug("new password: " .. topRec['password'])
-- return new password
return topRec['password']
end
```

Exercise 2 – Ruby: Register and Execute UDF

Locate the UserService module

1. In UserService.update_password
 1. Ensure the UDF module is registered
 2. Execute UDF

NOTE: UDF registration has been included here for convenience. The recommended way of registering UDFs in production environment is via AQL

In this exercise you will register and invoke the UDF created in Exercise 1.

We will programmatically register the UDF for convenience.

In UserService.update_password, locate these comments and add your code:

1. Ensure the UDF module is registered

```
task = client.register_udf_from_file(module_path,  
module_name, Language::LUA)  
task.wait_till_completed
```

2. Execute the UDF passing the new password
as a parameter, to the UDF

```
client.execute_udf(key, "updateUserPwd", "updatePassword", [new_password])
```

Summary

You have learned:

- Code a record UDF
- Register the UDF module
- Invoke a record UDF



Lab: Queries

Objectives

After successful completion of this Lab module you will have:

- Created a secondary index
- Prepared a statement
- Executed a query
- Processed the results

Lab Overview

The Lab exercises augment the "tweetaspike" application by allowing us to:

- 1) Query Tweets for a given username
- 2) Query users based on number of Tweets

The application shell is located in your cloned GitHub directory

`~/exercises/Queries/<language>`

Make sure you have your server up and you know its [IP address](#)

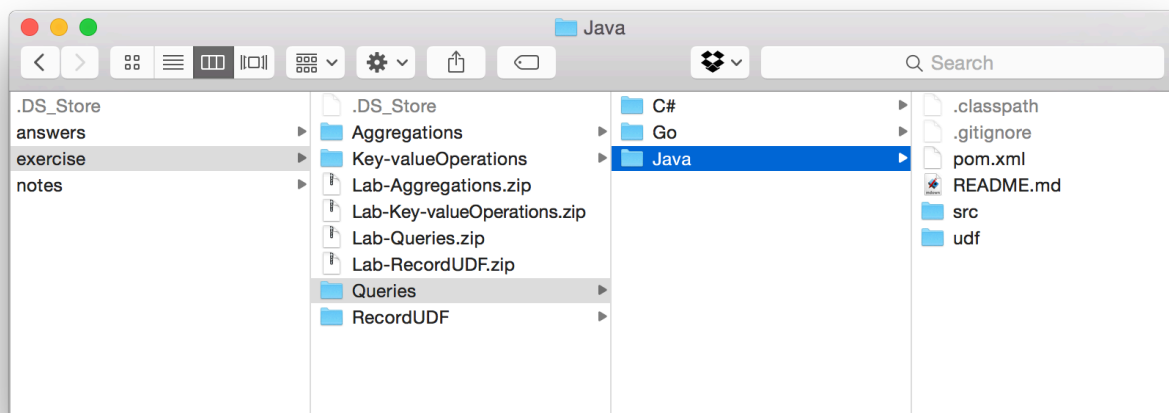
On your cloned or downloaded repository, you will find the following directories:

- Answers
- Exercise
- Notes

In the exercise directory, select the subdirectory for your programming language:

- C#
- Java
- Go
- Node.js
- PHP
- Python

The exercises for this module are in the Queries directory and you will find a Project/Solution/Codebase that is partly complete. Your task is to complete the code as outlined in each exercise.



Exercise 1 – Create secondary index on “tweetcount”

On your development cluster, create a secondary index using the **aql** utility:

1. Open a terminal connection to a node in your cluster
2. Execute the following AQL:

```
CREATE INDEX tweetcount_index ON test.users (tweetcount) NUMERIC
```
3. Verify the index status with the following AQL:

```
show indexes
```

Logon on to your server instance and run **aql** to create a numeric index on *tweetcount*.

At the prompt, enter the command:

```
CREATE INDEX tweetcount_index ON test.users (tweetcount) NUMERIC
```

Verify that the index has been created with the command:

```
show indexes
```

Exercise 2 – Create secondary index on “username”

On your development cluster, create a secondary index using the **aql** utility

1. Open a terminal connection to a node in your cluster
2. Execute the following AQL:

```
CREATE INDEX username_index ON test.tweets (username) STRING
```
3. Verify the index status with the following AQL:

```
show indexes
```

Logon on to your server instance and run **aql** to create a string index on *username*.

At the prompt, enter the command:

```
CREATE INDEX username_index ON test.tweets (username) STRING
```

Verify that the index has been created with the command:

```
show indexes
```



Ruby Exercises

Exercise 3 – Ruby: Query tweets for a given username

Locate the TweetService module in the Ruby project

In TweetService.query_tweets:

1. Optionally build a secondary index on the username of test.tweets
2. Create a query Statement
3. Create a filter predicate
4. Execute a query using:
 1. Namespace
 2. name of the set
 3. Filter for username

In TweetService.query_tweets add your code:

1. Create a Filter predicate
2. Execute a query using:
 1. the Namespace
 2. the Set name
 3. the Filter predicate to qualify the user name

```
statement = Statement.new('test', 'tweets', ['tweet'])
statement.filters << Filter.Equal('username', username)
results = client.query(statement)
```

Exercise 4 – Ruby: Query users based on num of tweets

Locate the module TweetService in the Rubyproject

In TweetService.query_by_tweetcount:

1. Create a range filter predicate for min-max range of tweetcount values.
2. Execute the query using:
 1. The namespace
 2. The set
 3. The range filter predicate
3. Execute query passing in null policy and instance of Statement created above
4. Iterate through returned RecordSet and for each record, output text in format
"<username> has <#> tweets"

In TweetService.query_by_tweet_count, add your code:

1. Create a range filter predicate for min-max tweetcount.
2. Execute the query using:
 1. The namespace
 2. The set
 3. The range filter predicate
3. Execute query passing in null policy and instance of Statement created above
4. Iterate through returned RecordSet and for each record, output text in format
"<username> has <#> tweets"

```
statement = Statement.new('test', 'users', ['username', 'tweetcount'])
statement.filters << Filter.Range('tweetcount', min, max) unless min == 0
&& max == 0
results = client.query(statement)
```

Summary

You have learned:

- How to create a secondary index
- How to create a Statement
- Execute a query on a secondary index
- Process the results from a query



Lab: Aggregations

Objective

After successful completion of this Lab module you will have:

- Coded a Stream UDF
- Register the UDF with a cluster
- Executed Aggregation from your C#, PHP, Node.js or Java application

Lab Overview

The lab exercise augments “tweetaspike” by using a Stream UDF. Here we will create a Stream UDF that aggregates number of users with tweet count between min-max range by region – North, South, East and West .

The application shell is located in your cloned GitHub directory
`~/exercises/Aggregations/<language>`

Make sure you have your server up and you know its [IP address](#)

In your cloned or downloaded repository, you will find the following directories:

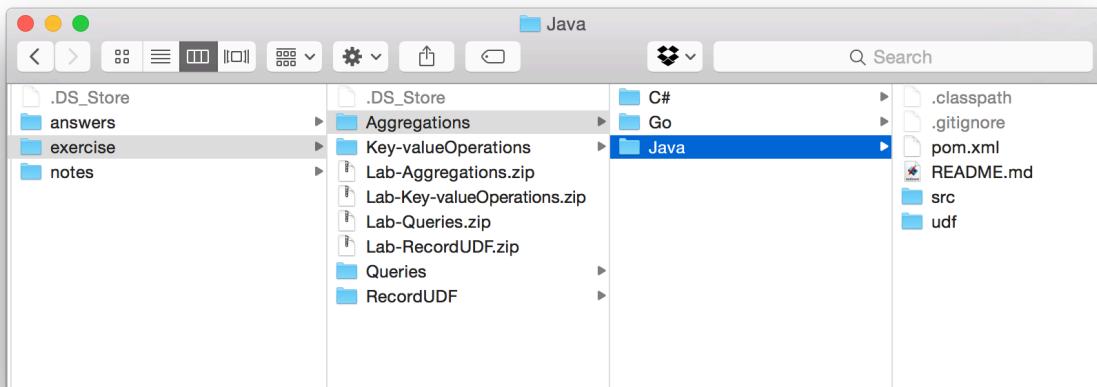
- Answers
- Exercise
- Notes

In the exercise directory, select the subdirectory for your programming language:

- C#
- Java
- Node.js
- PHP
- Python

The exercises for this module are in the Aggregations directory and you will find a Project/Solution/Codebase that is partly complete. Your task is to complete the code as outlined in each exercise.

Make sure you have your server up and you know its IP address



Exercise 1 – Write Stream UDF

Locate `aggregationByRegion.lua` file under `udf` folder in `AerospikeTraining` Solution

1. Code main function 'sum' to process incoming record stream and pass it to aggregate function 'aggregate_stats', then to reduce function 'reduce_stats'
2. Code aggregate function 'aggregate_stats' to examine value of 'region' bin and increment respective counters
3. Code reduce function 'reduce_stats' to merge maps

In this exercise you will create a Stream UDF module that:

- Aggregates (sums) tweets by region – The `aggregate_stats()` funcNon is invoked one for each element in the stream.
- Reduces the aggregations into a single Map of values – The `reduce_stats()` funcNon is invoked once for each data partition, once for each node in the cluster, and finally once on the client.
- The `sum()` funcNon configures the stream processing, and it is the funcNon invoked by the Client.

```
local function aggregate_stats(map,rec)
    -- Examine value of 'region' bin in record rec and increment respective counter in the map
    if rec.region == 'n' then
        map['n'] = map['n'] + 1
    elseif rec.region == 's' then
        map['s'] = map['s'] + 1
    elseif rec.region == 'e' then
        map['e'] = map['e'] + 1
    elseif rec.region == 'w' then
        map['w'] = map['w'] + 1
    end
    -- return updated map
    return map
end
local function reduce_stats(a,b)
    -- Merge values from map b into a
    a.n = a.n + b.n
    a.s = a.s + b.s
    a.e = a.e + b.e
    a.w = a.w + b.w
    -- Return updated map a
    return a
end
function sum(stream)
    -- Process incoming record stream and pass it to aggregate function, then to reduce function
    return stream : aggregate(map{n=0,s=0,e=0,w=0},aggregate_stats) : reduce(reduce_stats)
end
```

Summary

You have learned how to:

- Write a Stream UDF
- Write a Filter function
- Write a Map function
- Write an Aggregate function
- Write a Reduce function
- Execute an Aggregation from your application code



AEROSPIKE